

1 Q1

$$\begin{aligned}
L(z) &= \sum_{i=1}^n \|x^{(i)} - z\|^2 \\
\frac{\partial L}{\partial z_j} &= \sum_{i=1}^n \frac{\partial}{\partial z_j} \|x^{(i)} - z\|^2 = \sum_{i=1}^n \frac{\partial}{\partial z_j} (x_1^{(i)} - z_1)^2 + \dots + \frac{\partial}{\partial z_j} (x_j^{(i)} - z_j)^2 + \dots \\
&= \sum_{i=1}^n 0 + 2(x_j^{(i)} - z_j)(-1) + 0 = -2 \sum_{i=1}^n x_j^{(i)} - z_j \\
&= \nabla = \begin{bmatrix} \frac{\partial L}{\partial z_1} \\ \vdots \\ \frac{\partial L}{\partial z_d} \end{bmatrix} = -2 \sum_{i=1}^n x^{(i)} - z
\end{aligned}$$

Solve for gradient = 0 (and cancel the -2):

$$\begin{aligned}
\sum_{i=1}^n x^{(i)} - \sum_{i=1}^n z &= 0 \\
\sum_{i=1}^n x^{(i)} &= \sum_{i=1}^n z = nz \\
z &= \frac{\sum_{i=1}^n x^{(i)}}{n}
\end{aligned}$$

2 2

2.1 a

$$\begin{aligned}\frac{\partial L}{\partial w_j} &= \sum_{i=1}^n \frac{\partial}{\partial w_j} \left[(w \cdot x^{(i)}) \right] + \frac{\partial}{\partial w_j} \frac{1}{2} c \|w\|^2 \\ &= \sum_{i=1}^n x_j^{(i)} + c w_j \\ \nabla L(w) &= \begin{bmatrix} \frac{\partial L}{\partial w_1} \\ \vdots \\ \frac{\partial L}{\partial w_d} \end{bmatrix} = \sum_{i=1}^n x^{(i)} + c w\end{aligned}$$

2.2 b

$$\begin{aligned}\nabla L(w) &= \sum_{i=1}^n x^{(i)} + c w = 0 \\ c w &= - \sum_{i=1}^n x^{(i)} \\ w &= -\frac{1}{c} \sum_{i=1}^n x^{(i)}\end{aligned}$$

3 3

3.1 a

$$\nabla L(w) = \begin{bmatrix} \frac{\partial L}{\partial w_1} \\ \vdots \\ \frac{\partial L}{\partial w_4} \end{bmatrix} = \begin{bmatrix} 2w_1 + 2 \\ 4w_2 - 4 \\ 2w_3 - 2w_4 \\ 2w_4 - 2w_3 \end{bmatrix}$$

3.2 b

$$w_{t+1} = w_t - \eta \nabla L(w_t)$$

$$w_{t+1} = 0 - \eta \nabla L(0)$$

$$= \begin{bmatrix} 2 \\ -4 \\ 0 \\ 0 \end{bmatrix}$$

3.3 c

$$\nabla L(w) = \begin{bmatrix} \frac{\partial L}{\partial w_1} \\ \vdots \\ \frac{\partial L}{\partial w_4} \end{bmatrix} = \begin{bmatrix} 2w_1 + 2 \\ 4w_2 - 4 \\ 2w_3 - 2w_4 \\ 2w_4 - 2w_3 \end{bmatrix} = 0$$

$$w_1 = -1, w_2 = 1, w_3 = w_4$$

$$L((-1, 1, x, x)) = 1 + 2 + x^2 - 2x^2 + x^2 + 2(-1) - 4 + 4$$

$$L((-1, 1, x, x)) = 1$$

Minimum Value = 1

3.4 d

There is not a unique solution. A solution exists for all $w = (-1, 1, x, x)$.

4 4

4.1 a

$$\begin{aligned}
 \frac{\partial L}{\partial w_j} &= \sum_{i=1}^n \frac{\partial}{\partial w_j} \left[y^{(i)} - w \cdot x^{(i)} \right]^2 + \frac{\partial}{\partial w_j} \lambda \|w\|^2 \\
 &= \sum_{i=1}^n -2 \left[y^{(i)} - w \cdot x^{(i)} \right] x_j^{(i)} + 2\lambda w_j \\
 \nabla L(w) &= \begin{bmatrix} \frac{\partial L}{\partial w_1} \\ \vdots \\ \frac{\partial L}{\partial w_d} \end{bmatrix} = -2 \sum_{i=1}^n \left[y^{(i)} - w \cdot x^{(i)} \right] x^{(i)} + 2\lambda w
 \end{aligned}$$

4.2 b

$$\begin{aligned}
 w_{t+1} &= w_t - \eta_t \nabla L(w_t) \\
 w_{t+1} &= w_t - \eta_t \left\{ -2 \sum_{i=1}^n \left[y^{(i)} - w \cdot x^{(i)} \right] x^{(i)} + 2\lambda w \right\}
 \end{aligned}$$

4.3 c

Update Equation:

$$\begin{aligned}
 w_{t+1} &= w_t - \eta_t \nabla l(w_t; x^{(i)}, y^{(i)}) \\
 w_{t+1} &= w_t - \eta_t \left\{ -2 \left[y^{(i)} - w \cdot x^{(i)} \right] x^{(i)} + 2\lambda w \right\}
 \end{aligned}$$

Algorithm:

$w_0 = 0$

Cycle through points $(x^{(i)}, y^{(i)})$ until some stopping condition:

- $w_{t+1} = w_t - \eta_t \left\{ -2 \left[y^{(i)} - w_t \cdot x^{(i)} \right] x^{(i)} + 2\lambda w_t \right\}$

5 5

5.1 a

$$\nabla^2 f(x) = 2 \geq 0, \text{ convex.}$$

5.2 b

$$\nabla^2 f(x) = -2 \leq 0, \text{ concave.}$$

5.3 c

$$\nabla^2 f(x) = 2 \geq 0, \text{ convex.}$$

5.4 d

$$\nabla^2 f(x) = 0 = 0, \text{ both.}$$

5.5 e

$$\nabla^2 f(x) = 6x, \text{ hessian can change signs, neither.}$$

5.6 f

$$\nabla^2 f(x) = 12x^2 \geq 0, \text{ convex.}$$

5.7 g

$$\nabla f(x) = \frac{1}{x} = x^{-1}, \nabla^2 f(x) = -x^{-2} \leq 0, \text{ concave.}$$

6 6

6.1 a

For my coordinate descent method, I will repeatedly call a routine until the gradient magnitude is less than a certain threshold. For this routine:

- (1) Initialize the weight vector to a d-dimensional vector sampled from a standard normal (seeded for reproducibility)
- (2) while the gradient magnitude is greater than a certain value:
 - (3) Compute the gradient
 - (4) (i) Pick the coordinate with the largest absolute value. This, intuitively, is the component-direction of steepest ascent (and thus the negative is the direction of steepest descent)
 - (5) (ii) Update this coordinate by the magnitude of the gradient. Thus, when the model is close to optimized, the step size will naturally decrease in the corresponding direction.
 $w[\text{idx}] = w[\text{idx}] - \eta * \text{sign}(\nabla[\text{idx}]) * \text{norm}(\nabla)$

Because this method depends on the gradient, the function must be defined and differentiable at every point in R^d .

6.2 b

In [149...

```
import pandas as pd
import numpy as np
import os
from pathlib import Path
from sklearn import linear_model, metrics

ITERS=100000

df = pd.read_csv("data/HW4/heart.csv")
df.head()

X = np.array(df.iloc[:, :-1])
y = np.array(df.iloc[:, [-1]])

np.random.seed(123)
idx = np.random.choice(303, 303, replace=False)
tr_idx, te_idx = idx[:200], idx[200:]

X_tr_raw, y_tr_raw = X[tr_idx], y[tr_idx]
X_te_raw, y_te_raw = X[te_idx], y[te_idx]

def normalize(mode, X_tr, X_te):
    if mode == "z":
        # z-score normalize
        X_tr_mu = np.mean(X_tr, axis=0, keepdims=True)
        X_tr_sig = np.std(X_tr, axis=0, keepdims=True)

        X_tr_ = (X_tr - X_tr_mu) / X_tr_sig
        X_te_ = (X_te - X_tr_mu) / X_tr_sig
        return X_tr_, X_te_
    elif mode == "min-max":
        # min-max 0-1 normalization
        X_tr_mi = np.min(X_tr, axis=0, keepdims=True)
        X_tr_ma = np.max(X_tr, axis=0, keepdims=True)
        X_tr_ = (X_tr - X_tr_mi) / (X_tr_ma - X_tr_mi)
        X_te_ = (X_te - X_tr_mi) / (X_tr_ma - X_tr_mi)
        return X_tr_, X_te_
    else:
        return X_tr, X_te

def logit(w, b, X):
    return X @ w + b

def logit_nob(w, X):
    return X @ w

def prob(w, b, X, y):
    logits = logit(w, b, X)
    assert(y.shape == logits.shape)
    return 1 / (1 + np.exp(-y * logits))

def prob_nob(w, X, y):
    logits = logit_nob(w, X)
```



```

    assert(y.shape == logits.shape)
    return 1 / (1 + np.exp(-y * logits))

def loss(w, b, X, y):
    logits = logit(w, b, X)
    assert(y.shape == logits.shape)
    return np.mean(np.log(1 + np.exp(-y * logits)))

def loss_nob(w, X, y):
    logits = logit_nob(w, X)
    assert(y.shape == logits.shape)
    return np.mean(np.log(1 + np.exp(-y * logits)))

def test_err(reg: linear_model.LogisticRegression, X, y):
    y_hat = reg.predict(X)
    err = np.not_equal(y_hat, y.flatten())
    return np.sum(err)

def reg_loss(reg: linear_model.LogisticRegression, X, y):
    w, b = reg.coef_.reshape((-1, 1)), reg.intercept_
    wx_b = reg.decision_function(X).reshape((-1, 1))
    assert(np.allclose(wx_b, logit(w, b, X)))
    return loss(w, b, X, y)

def experiment(mode, X_tr, X_te, y_tr, y_te):
    print(f"=== Using {mode} Normalization ===")
    X_tr_, X_te_ = normalize(mode, X_tr, X_te)
    reg = linear_model.LogisticRegression(C=999999999).fit(X_tr_, y_tr.flatten())
    w, b = reg.coef_.reshape((-1, 1)), reg.intercept_

    print("\tw: ", np.round(w, 4).flatten())
    print("\tb: ", np.round(b, 4))

    # print(f"=== Testing ===")
    # err = test_err(reg, X_te_, y_te)
    # print(f"\tError Rate: {np.sum(err)}/{len(y_te)}={np.sum(err)/len(y_te):0.4f}")
    y_1s_tr = np.where(y_tr == 0, -1, 1)
    y_1s_te = np.where(y_te == 0, -1, 1)

    tr_loss = reg_loss(reg, X_tr_, y_1s_tr)
    te_loss = reg_loss(reg, X_te_, y_1s_te)
    print(f"\tTrain Loss: {tr_loss}, Test Loss: {te_loss}")

    # print(y_tr[:10], y_1s_tr[:10])
    # y_hat = prob(w, b, X_tr_, y_1s_tr)
    # y_hat2 = reg.predict_proba(X_tr_)
    # print(y_hat[:10], y_hat2[:10])
    # calc_loss = metrics.log_loss(y_tr, y_hat2[:,1].flatten())
    # print(f"\tCalculated Loss: {calc_loss}")
    print()
    print("-" * 50)
    print()
    return tr_loss, te_loss

```

In [150...

```
tr_loss, te_loss = experiment(  
    "min-max", X_tr_raw, X_te_raw, y_tr_raw, y_te_raw)  
  
=== Using min-max Normalization ===  
    w: [-0.0916 -2.3859  2.5735 -3.329  -3.3522  0.5152  0.6446  4.345  -0.9711  
-4.1391  1.8231 -3.505  -2.4696]  
    b: [2.4201]  
    Train Loss: 0.3138296347754303, Test Loss: 0.44660770010026035  
  
-----
```

In [154...

```
import tqdm

def compute_gradient(w, X, y):
    prob_correct = prob_nob(w, X, y)
    prob_wrong = 1 - prob_correct
    assert(y.shape == prob_wrong.shape)
    matrix = y * X * prob_wrong
    return -np.mean(matrix, axis=0).reshape((-1, 1))

def experiment_coordinate(mode, thresh, lr, X_tr, y_tr, X_te, y_te, sel_coord="mag", it
    print(
        f"*** Using {mode} Normalization, thresh {thresh:0.5f}, Coordinate Selector mode
    np.random.seed(123)
    X_tr_, X_te_ = normalize(mode, X_tr, X_te)
    n_tr, n_te = X_tr_.shape[0], X_te_.shape[0]

    # add bias to data
    X_tr_, X_te_ = np.hstack(
        (X_tr_, np.ones((n_tr, 1)))
    ), np.hstack((X_te_, np.ones((n_te, 1))))

    w = np.random.normal(size=(X_tr_.shape[1], 1))
    losses = [loss_nob(w, X_tr_, y_tr)]
    te_losses = [loss_nob(w, X_te_, y_te)]

    for i in range(iters):
        gradient = compute_gradient(w, X_tr_, y_tr)
        gradient_norm = np.linalg.norm(gradient.flatten())
        if gradient_norm <= thresh:
            # print(f"Breaking on iteration [{i}]")
            break

        if sel_coord == "mag":
            idx_big = np.argmax(np.abs(gradient).flatten())
        elif sel_coord == "random":
            idx_big = np.random.randint(0, len(w), size=1).item()

        sign = np.sign(gradient[idx_big, 0])
        w[idx_big, 0] -= lr * sign * gradient_norm
        losses.append(loss_nob(w, X_tr_, y_tr))
        te_losses.append(loss_nob(w, X_te_, y_te))
    print("-" * 50)

    return w, losses, te_losses, X_tr_, X_te_
```

In [155...

```
import matplotlib.pyplot as plt

def batch_experiment(thresh=1e-1, lr=1e-3, iters=10000):
    _, losses_mag, te_losses_mag, _, _ = experiment_coordinate(
        "min-max", thresh, lr, X_tr_raw, y_tr_raw, X_te_raw, y_te_raw, sel_coord="mag", i
    )
    _, losses_rdm, te_losses_rdm, _, _ = experiment_coordinate(
        "min-max", thresh, lr, X_tr_raw, y_tr_raw, X_te_raw, y_te_raw, sel_coord="random"
    )
```

```

ct = max(len(losses_mag), len(losses_rdm))
x_ax = np.arange(0, ct)
line = np.ones_like(x_ax)
plt.plot(range(len(losses_mag)), losses_mag, "bo-", markevery=1000, label="Coordinate
# plt.plot(x_ax, te_losses_mag, "co-", markevery=1000, label="Coordinate Descent (min
plt.plot(x_ax, line * tr_loss, "r-", label="sklearn (train)")
# plt.plot(x_ax, line * te_loss, "y-", label="sklearn (test)")
plt.plot(range(len(losses_rdm)), losses_rdm, "go-", label="Coordinate Descent (random
# plt.plot(x_ax, te_losses_rdm, "mo-", label="Coordinate Descent (random) (te)", mark
plt.legend()
plt.title(f"Loss Plots ($\lambda$={lr},thresh={thresh}). $L(w) = \frac{1}{n} \sum$
plt.show()

```

In [156...

```

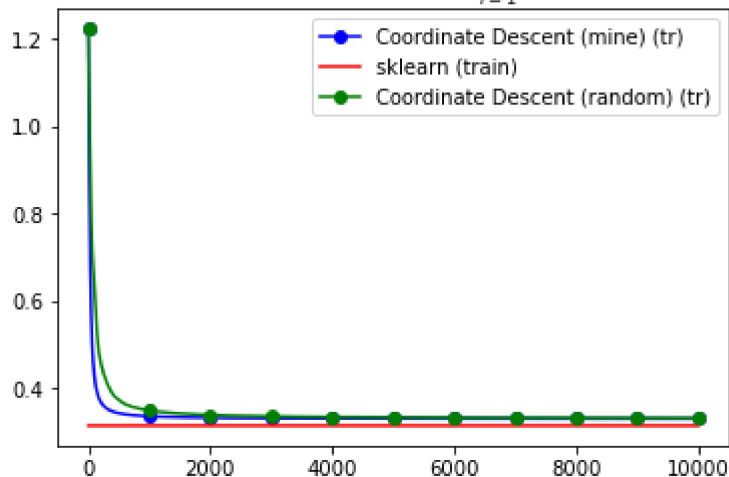
batch_experiment(thresh=1e-3,lr=1e-1)
batch_experiment(thresh=1e-3,lr=1e-2)
batch_experiment(thresh=1e-3,lr=1e-3)
batch_experiment(thresh=1e-3,lr=1e-4)
batch_experiment(thresh=1e-3,lr=1e-5)

```

*** Using min-max Normalization, thresh 0.00100, Coordinate Selector mode mag

*** Using min-max Normalization, thresh 0.00100, Coordinate Selector mode random

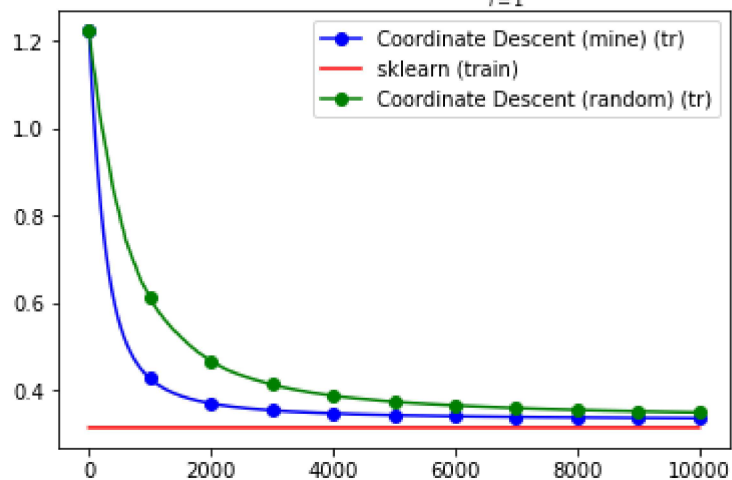
Loss Plots ($\lambda=0.1$,thresh=0.001). $L(w) = \frac{1}{n} \sum_{i=1}^n \ln(1 + \exp(-y^{(i)}(w \cdot x^{(i)})))$



*** Using min-max Normalization, thresh 0.00100, Coordinate Selector mode mag

*** Using min-max Normalization, thresh 0.00100, Coordinate Selector mode random


Loss Plots ($\lambda=0.01, \text{thresh}=0.001$). $L(w) = \frac{1}{n} \sum_{i=1}^n \ln(1 + \exp(-y^{(i)}(w \cdot x^{(i)})))$



*** Using min-max Normalization, thresh 0.00100, Coordinate Selector mode mag

*** Using min-max Normalization, thresh 0.00100, Coordinate Selector mode random


Loss Plots ($\lambda=0.001, \text{thresh}=0.001$). $L(w) = \frac{1}{n} \sum_{i=1}^n \ln(1 + \exp(-y^{(i)}(w \cdot x^{(i)})))$



*** Using min-max Normalization, thresh 0.00100, Coordinate Selector mode mag

*** Using min-max Normalization, thresh 0.00100, Coordinate Selector mode random

Loss Plots ($\lambda=0.0001, \text{thresh}=0.001$). $L(w) = \frac{1}{n} \sum_{i=1}^n \ln(1 + \exp(-y^{(i)}(w \cdot x^{(i)})))$



*** Using min-max Normalization, thresh 0.00100, Coordinate Selector mode mag

*** Using min-max Normalization, thresh 0.00100, Coordinate Selector mode random

Loss Plots ($\lambda=1e-05, \text{thresh}=0.001$). $L(w) = \frac{1}{n} \sum_{i=1}^n \ln(1 + \exp(-y^{(i)}(w \cdot x^{(i)})))$

