

James Zhao
HW5
A15939512

1 1

$$\nabla F(x) = 2(x - u)$$

$$\nabla^2 F(x) = 2I$$

A diagonal matrix is positive semidefinite, thus it is convex.

2 2

The function is a sum of terms, thus each gradient element is only a function of p_i . This means all basic operations are vectorizable.

$$\nabla F(p) = -(p * 1/p + \ln(p)) = -(1 + \ln(p))$$

$$\nabla^2 F(p) = \text{diag}(-1/p)$$

The hessian is a diagonal matrix of negative values (since p ranges from 0 to 1). Thus, entropy must be concave.

3 3

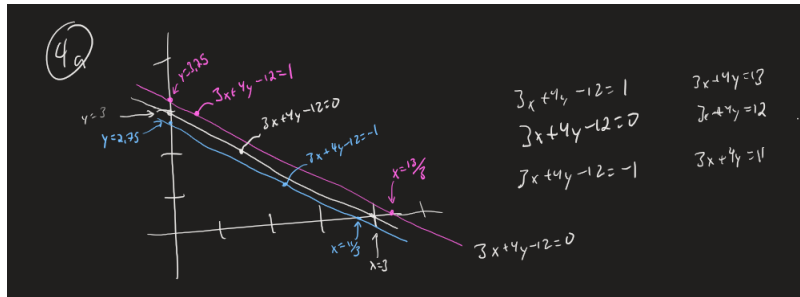
In the perceptron learning algorithm, b is updated by the label everytime an update is performed. Thus,

$$b_0 = 0, b_{final} = 0 + (-1) * p + (1) * q$$

$$b_{final} = q - p$$

4 4

4.1 a



4.2 b

4.3 c

$$\gamma = \frac{1}{\|w\|} = \frac{1}{5}$$

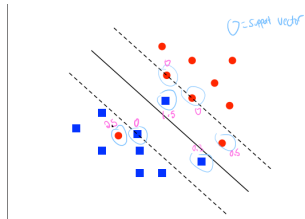
4.4 d

$$\begin{aligned} \text{classification} &= \text{sign}(w \cdot x + b) \\ &= \text{sign}(3 * 2 + 4 * 2 - 12) = \text{sign}(2) = +1 \end{aligned}$$

5 5

5.1 a

5.2 b



If the C parameter increased, we would be punished more for using slack. Thus, the margin would grow, since the magnitude of slack is in terms of the margin size (same absolute error / larger margin = smaller slack).

6 6

6.1 a

Possibly false. The α_i are incremented every time a point is classified wrong, and the point can be classified wrong multiple times.

6.2 b

Necessarily true. Each update increments one entry in the α vector. Thus, the sum of entries must equal the number of updates.

6.3 c

Necessarily true. If we wanted to maximize the number of nonzero entries in the α vector, we can update a different entry each time. This leads to at most k indices having non-zero entries.

6.4 d

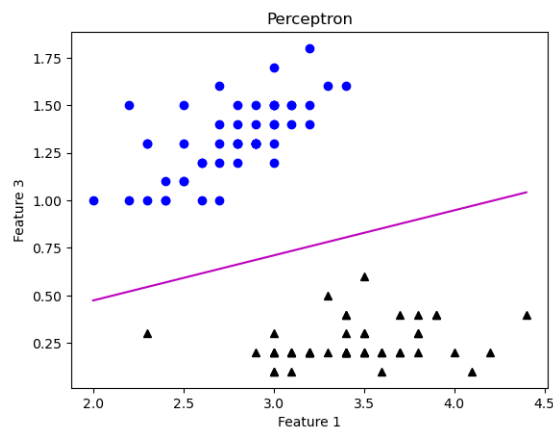
Necessarily true. If a perceptron converges, it means that there are no more possible updates / no misclassified points and the binary classifier is perfectly modeled. This means that there must exist a linear decision boundary - hence the data must be linearly separable.

7 7

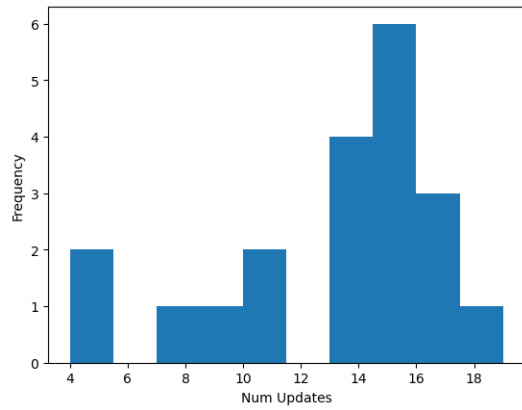
7.1 a

7.2 b

7.3 c



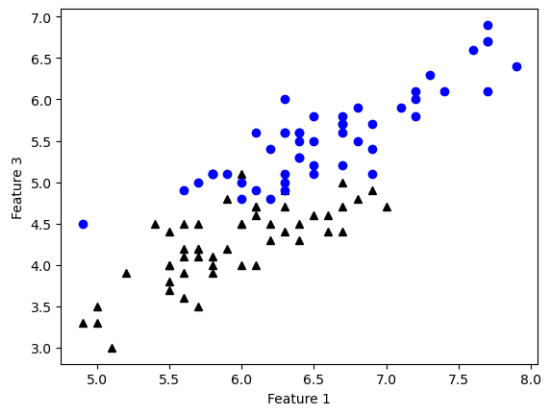
7.4 d



8 8

8.1 a

No.




8.2 b

| C | Error | Num Support Vectors |
|-------|-------|---------------------|
| 0.125 | 0.07 | 52 |
| 0.25 | 0.06 | 45 |
| 0.5 | 0.06 | 38 |
| 1.0 | 0.07 | 31 |
| 2.0 | 0.06 | 24 |
| 4.0 | 0.07 | 21 |
| 8.0 | 0.05 | 19 |
| 16.0 | 0.07 | 16 |
| 32.0 | 0.06 | 15 |
| 64.0 | 0.05 | 14 |
| 128.0 | 0.05 | 14 |
| 256.0 | 0.05 | 14 |
| 512.0 | 0.05 | 14 |

8.3 c

Best value: $C=8$



```
In [96]: import numpy as np
```

Q7

```
In [97]: # (7a)
# First Function
def perceptron(w, b, x):
    return np.sign(np.dot(w.flatten(), x.flatten()) + b)

def has_converged(w, b, X, l):
    return np.allclose(np.sign(X @ w + b).flatten(), l.flatten())

MAX_ITERS = 72

# Second Function
def train(X, l):
    # data: n x d, labels: n
    n, d = X.shape
    assert(l.shape == (n,))

    w = np.zeros((d, 1))
    b = 0

    num_updates = 0

    for i in range(MAX_ITERS):
        permutation = np.random.choice(range(n), n, replace=False)
        X_perm = X[permutation]
        l_perm = l[permutation]

        for j in range(n):
            x_j = X_perm[[j],:]
            l_j = l_perm[j]
            if perceptron(w, b, x_j) != l_j:
                num_updates += 1
                w = w + l_j * x_j.reshape((-1, 1))
                b += l_j

        if has_converged(w, b, X_perm, l_perm):
            break

    return w, b, num_updates
```

```
In [98]: # (7b, 7c)

from sklearn import datasets
iris = datasets.load_iris()
X0 = iris.data
y0 = iris.target

print(X0.shape, y0.shape)

rows = np.array(np.nonzero(y0 <= 1)[0])

X_1 = X0[rows.reshape((-1, 1)),(1,3)]
```

```
y_1 = np.where(y0[rows] == 0, -1, 1)
```

```
print(X_1.shape, y_1.shape)
```

```
(150, 4) (150,)
```

```
(100, 2) (100,)
```

In [99]:


```
import matplotlib.pyplot as plt

# (7c)
np.random.seed(123)
w, b, num_updates = train(X_1, y_1)

def plot(w, b, X, l, draw_line = True, title = "Perceptron"):
    w1, w2 = w.flatten()
    x_min, x_max = np.min(X[:,0]), np.max(X[:,0])
    line_x = np.linspace(x_min, x_max, 100)
    #  $w_1x + w_2y + b = 0$ 
    #  $y = -w_1/w_2 x - b / w_2$ 
    line_y = - w1 / w2 * line_x - b / w2

    plt.plot(X[l==0], X[l==1], "^k")
    plt.plot(X[l==1,0], X[l==1,1], "ob")
    if draw_line:
        plt.plot(line_x, line_y, "-m")
    plt.xlabel("Feature 1")
    plt.ylabel("Feature 3")
    plt.title(title)


plot(w, b, X_1, y_1)
```

In [100...

```
np.random.seed(1337)
n_trials = 20
updates = []
for i in range(n_trials):
    _, _, n_i = train(X_1, y_1)
    updates.append(n_i)
total_updates = sum(updates)
print(f"Average number of updates: {total_updates}/{n_trials}={total_updates / n_trials}")
plt.hist(updates)
plt.xlabel("Num Updates")
plt.ylabel("Frequency");
```

Average number of updates: 258/20=12.9



Q8

In [101...


```
rows_2 = np.array(np.nonzero(y0 > 0)[0])
X_2 = X0[rows_2.reshape((-1, 1)), (0, 2)]
y_2 = np.where(y0[rows_2] == 1, -1, 1)

print(X0.shape, y0.shape)
print(X_2.shape, y_2.shape)

plot(np.ones((2, 1)), 1, X_2, y_2, draw_line = False, title = None)

print("(8a) The data is NOT linearly seperable")
```

```
(150, 4) (150,)
(100, 2) (100,)
(8a) The data is NOT linearly seperable
```



```
In [102... np.c_[np.array([1,2,3]), np.array([4,5,6])]
```

```
Out[102... array([[1, 4],
        [2, 5],
        [3, 6]])
```

```
In [103... import sklearn
def full_plot(X, y, svc: sklearn.svm.SVC, C, error):
    plt.plot(X[y==-1,0], X[y==-1,1], "^k", markersize=3)
    plt.plot(X[y==1,0], X[y==1,1], "ob", markersize=3)
    margin = 0.25
    x_min, x_max = np.min(X[:,0]), np.max(X[:,0])
    y_min, y_max = np.min(X[:,1]), np.max(X[:,1])
    delta = 0.01

    xx, yy = np.meshgrid(np.arange(x_min - margin, x_max + margin, delta), np.arange(y_min - margin, y_max + margin, delta))
    Z = svc.decision_function(np.c_[xx.flatten(), yy.flatten()])
    for i in range(len(Z)):
        Z[i] = min(Z[i], 1.0)
        Z[i] = max(Z[i], -1.0)
        if (Z[i] > 0.0) and (Z[i] < 1.0):
            Z[i] = 0.5
        if (Z[i] < 0.0) and (Z[i] > -1.0):
            Z[i] = -0.5
    Z = Z.reshape(xx.shape)
    plt.pcolormesh(xx, yy, Z, cmap=plt.cm.BuPu, vmin=-2, vmax=2)
    plt.xlim((x_min - margin, x_max + margin))
    plt.ylim((y_min - margin, y_max + margin))
    plt.title(f"C={C}, Error={error}")
    plt.show()
```

In [104...


```
import sklearn

Cs = np.exp2(np.arange(-3, 10))
data = []
for c in Cs:
    print(f"=== {c} ===")
    svm = sklearn.svm.SVC(kernel="linear", C=c)
    svc = svm.fit(X=X_2, y=y_2)


    wrong = np.mean(np.not_equal(svm.predict(X_2), y_2))
    print(f"\tError rate: {wrong}")
    full_plot(X_2, y_2, svc, c, wrong)
    num_sv = len(svc.support_)

    data.append((c, wrong, num_sv))
```


```
=== 0.125 ===
    Error rate: 0.07
```




```
=== 0.25 ===
    Error rate: 0.06
```




=== 0.5 ===
Error rate: 0.06




=== 1.0 ===
Error rate: 0.07




=== 2.0 ===
Error rate: 0.06




=== 4.0 ===
Error rate: 0.07




=== 8.0 ===
Error rate: 0.05




=== 16.0 ===
Error rate: 0.07




=== 32.0 ===
Error rate: 0.06




=== 64.0 ===
Error rate: 0.05




=== 128.0 ===
Error rate: 0.05



=== 256.0 ===
Error rate: 0.05



=== 512.0 ===
Error rate: 0.05



In [105...

```
for v in data:
    print(str(v)[1:-1].replace(',', ' &') + "\\\\\\\hline")
```

0.125 & 0.07 & 52\\\hline
0.25 & 0.06 & 45\\\hline
0.5 & 0.06 & 38\\\hline
1.0 & 0.07 & 31\\\hline
2.0 & 0.06 & 24\\\hline
4.0 & 0.07 & 21\\\hline
8.0 & 0.05 & 19\\\hline
16.0 & 0.07 & 16\\\hline
32.0 & 0.06 & 15\\\hline
64.0 & 0.05 & 14\\\hline
128.0 & 0.05 & 14\\\hline
256.0 & 0.05 & 14\\\hline
512.0 & 0.05 & 14\\\hline